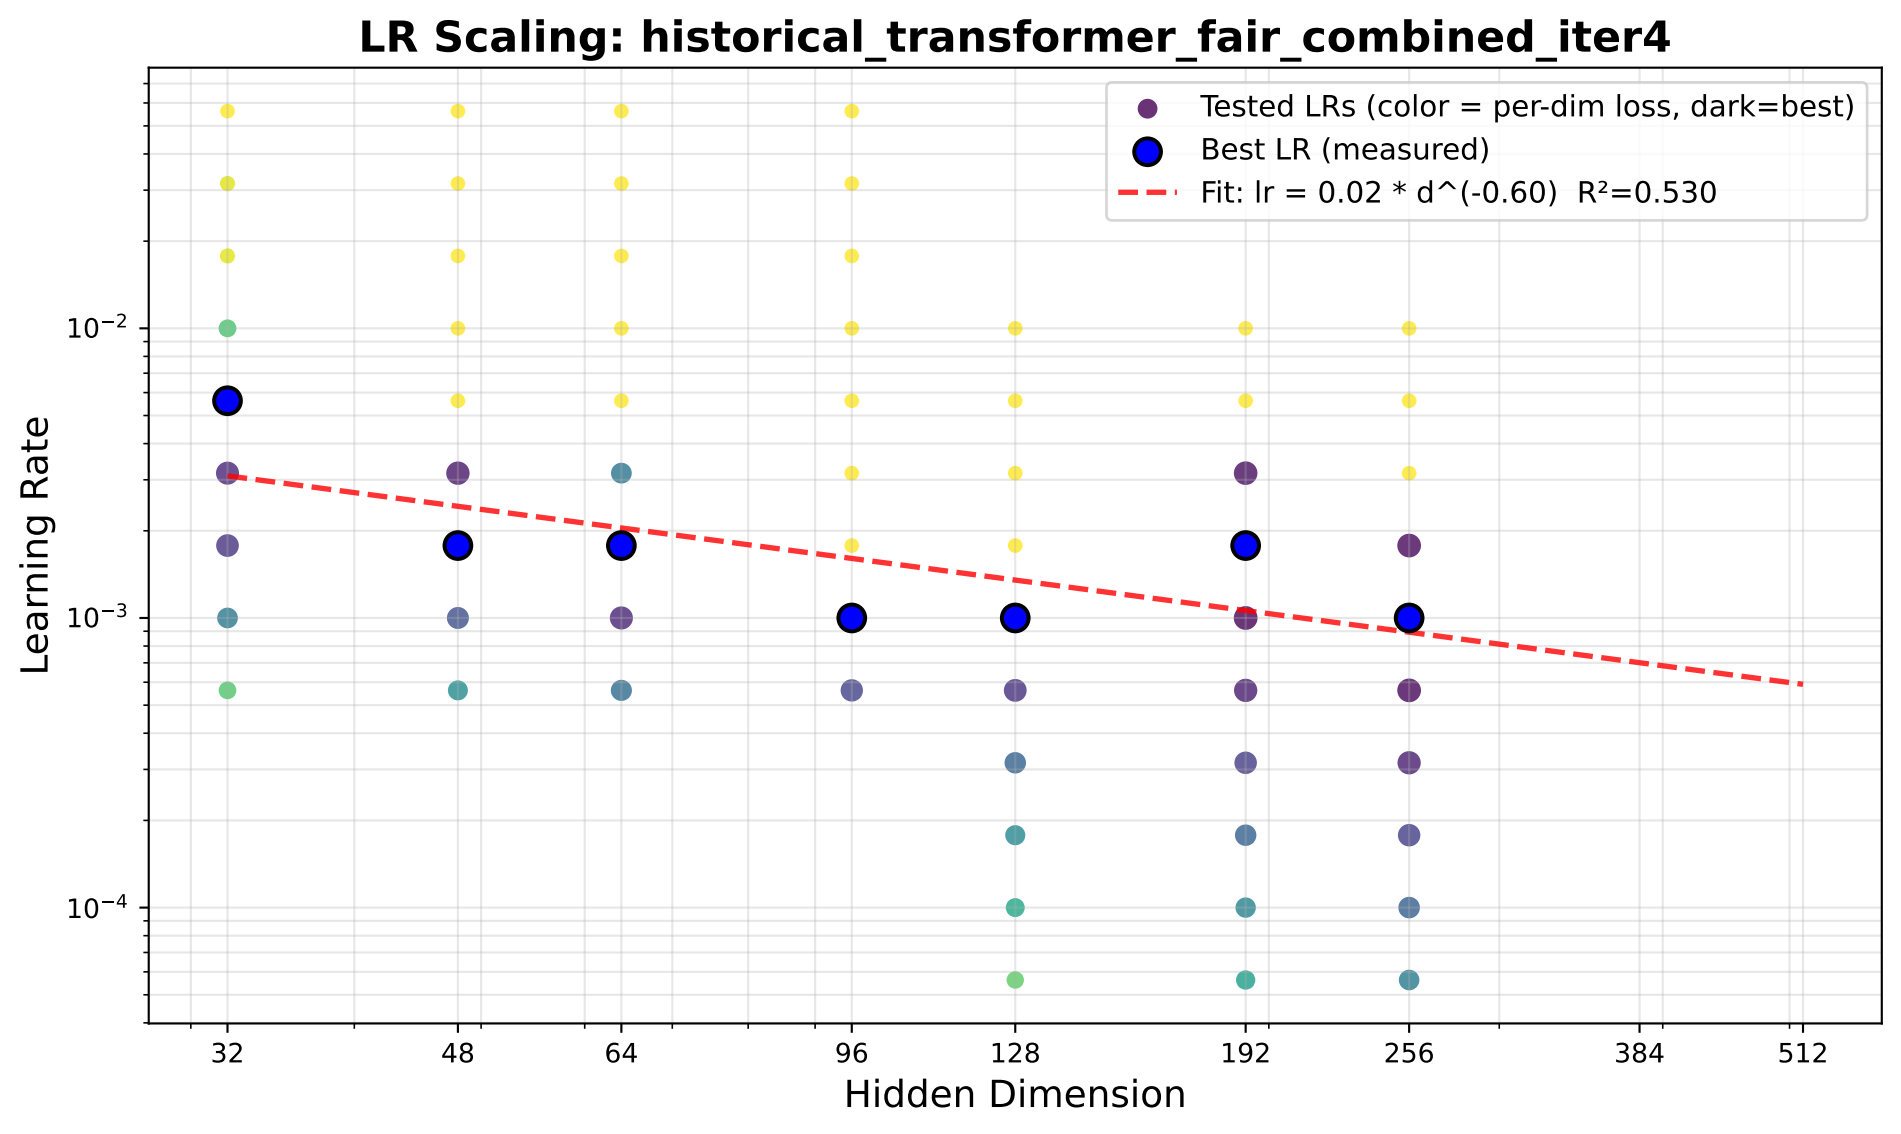
LR Scaling: historical_transformer_fair_combined_iter4
Tested LRs (color = per-dim loss, dark=best)
Best LR (measured)
Fit: lr = 0.02 * d^(-0.60) R²=0.530
Learning Rate
10^{-2}
10^{-3}
10^{-4}
32
48
64
96
128
192
256
384
512
Hidden Dimension