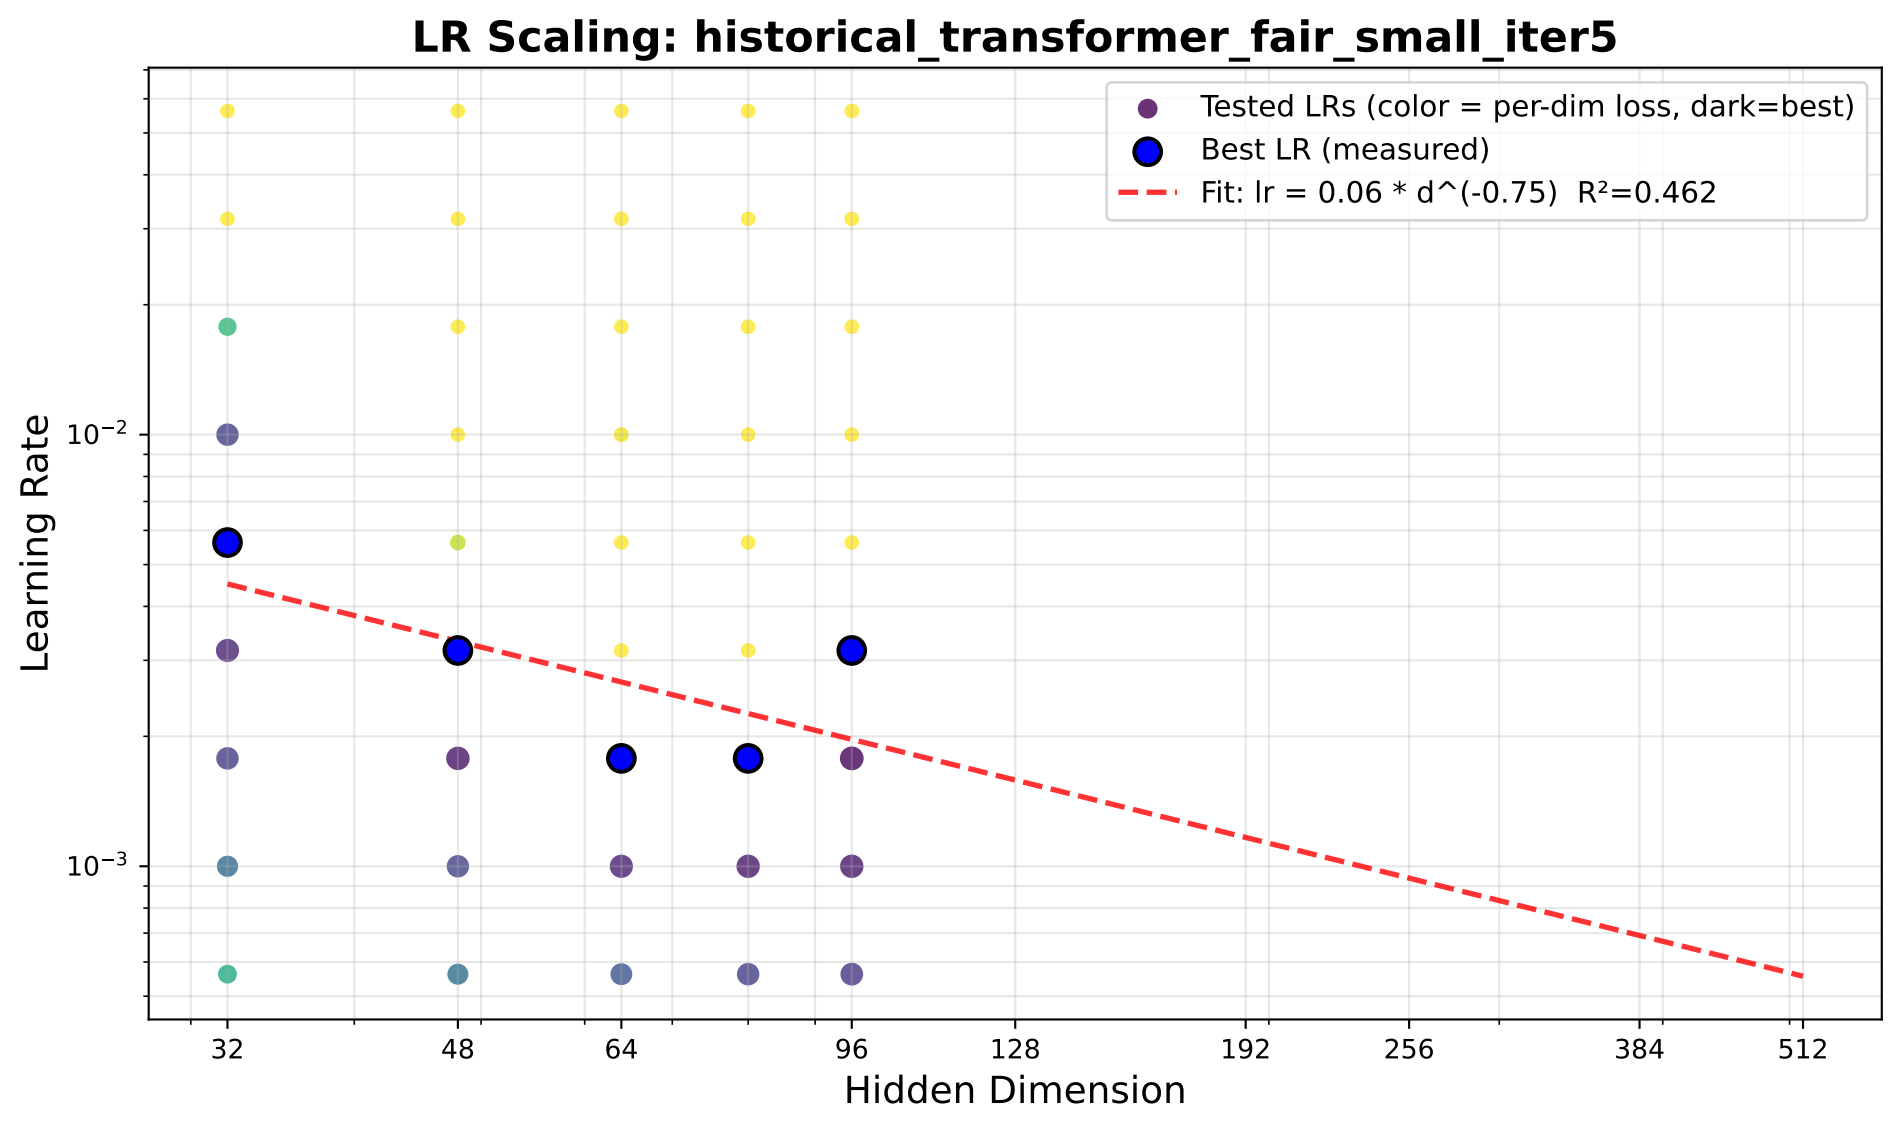
LR Scaling: historical_transformer_fair_small_iter5
Tested LRs (color = per-dim loss, dark=best)
Best LR (measured)
Fit: lr = 0.06 * d^(-0.75) R²=0.462
Learning Rate
10^{-2}
10^{-3}
32
48
64
96
128
192
256
384
512
Hidden Dimension