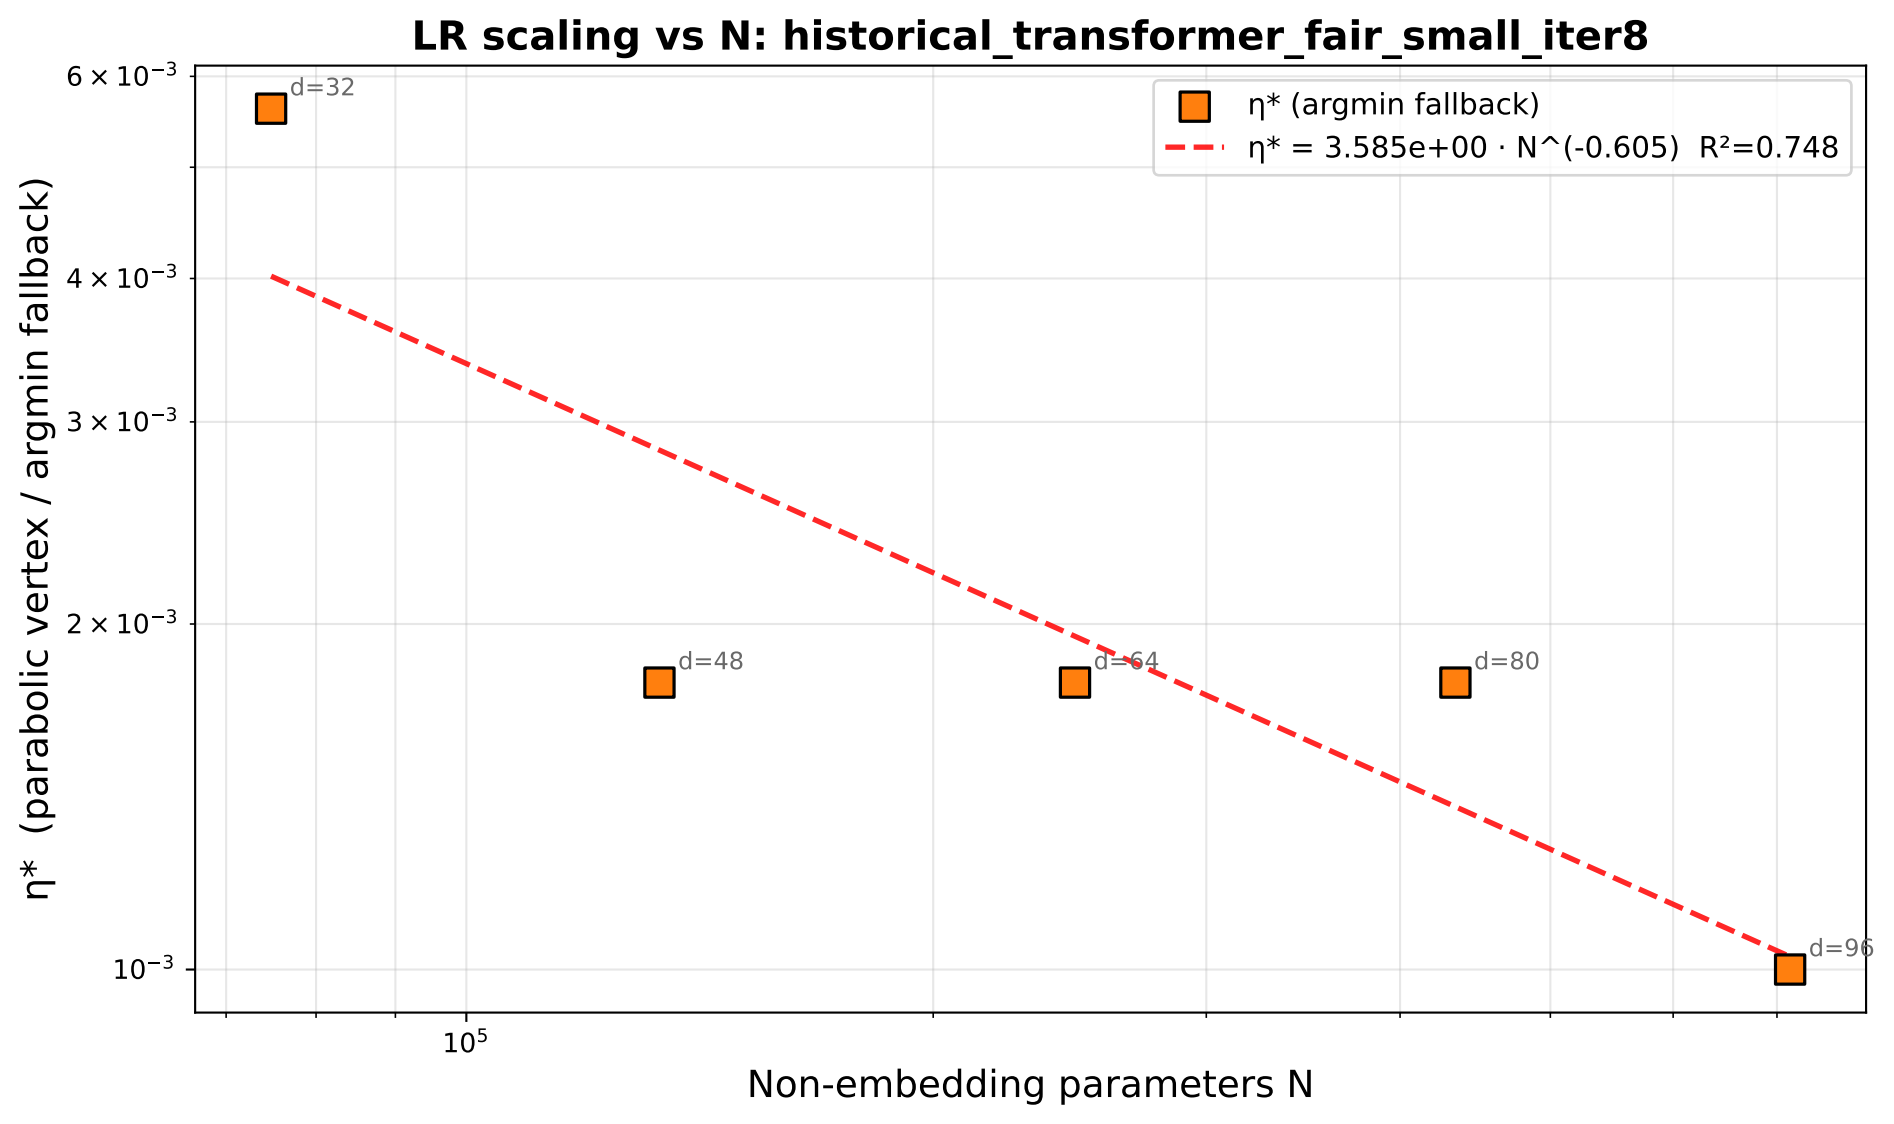
LR scaling vs N: historical_transformer_fair_small_iter8
η* (argmin fallback)
η* = 3.585e+00 · N^(-0.605) R²=0.748
d=32
d=48
d=64
d=80
d=96
6 × 10⁻³
4 × 10⁻³
3 × 10⁻³
2 × 10⁻³
10⁻³
10⁵
η* (parabolic vertex / argmin fallback)
Non-embedding parameters N