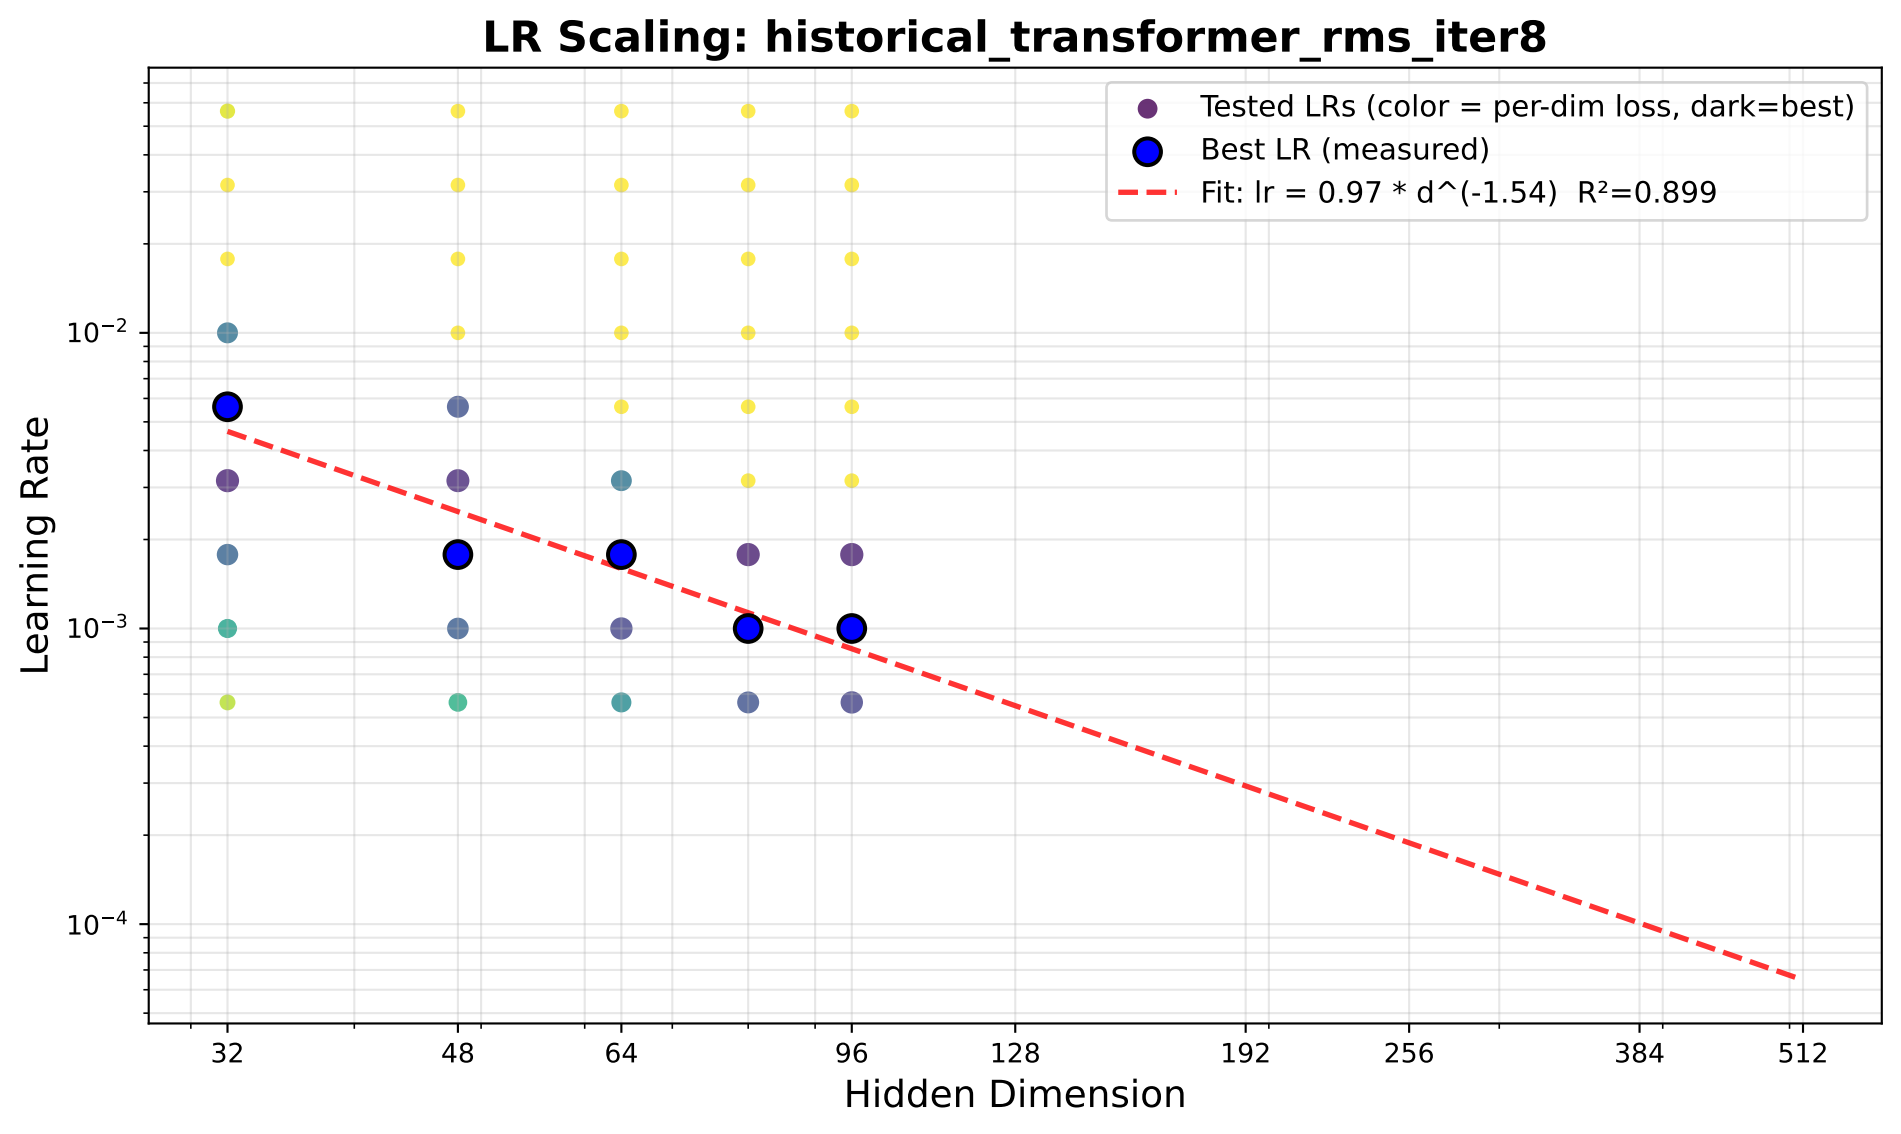
LR Scaling: historical_transformer_rms_iter8
Tested LRs (color = per-dim loss, dark=best)
Best LR (measured)
Fit: lr = 0.97 * d^(-1.54) R²=0.899
Learning Rate
10^{-2}
10^{-3}
10^{-4}
32
48
64
96
128
192
256
384
512
Hidden Dimension