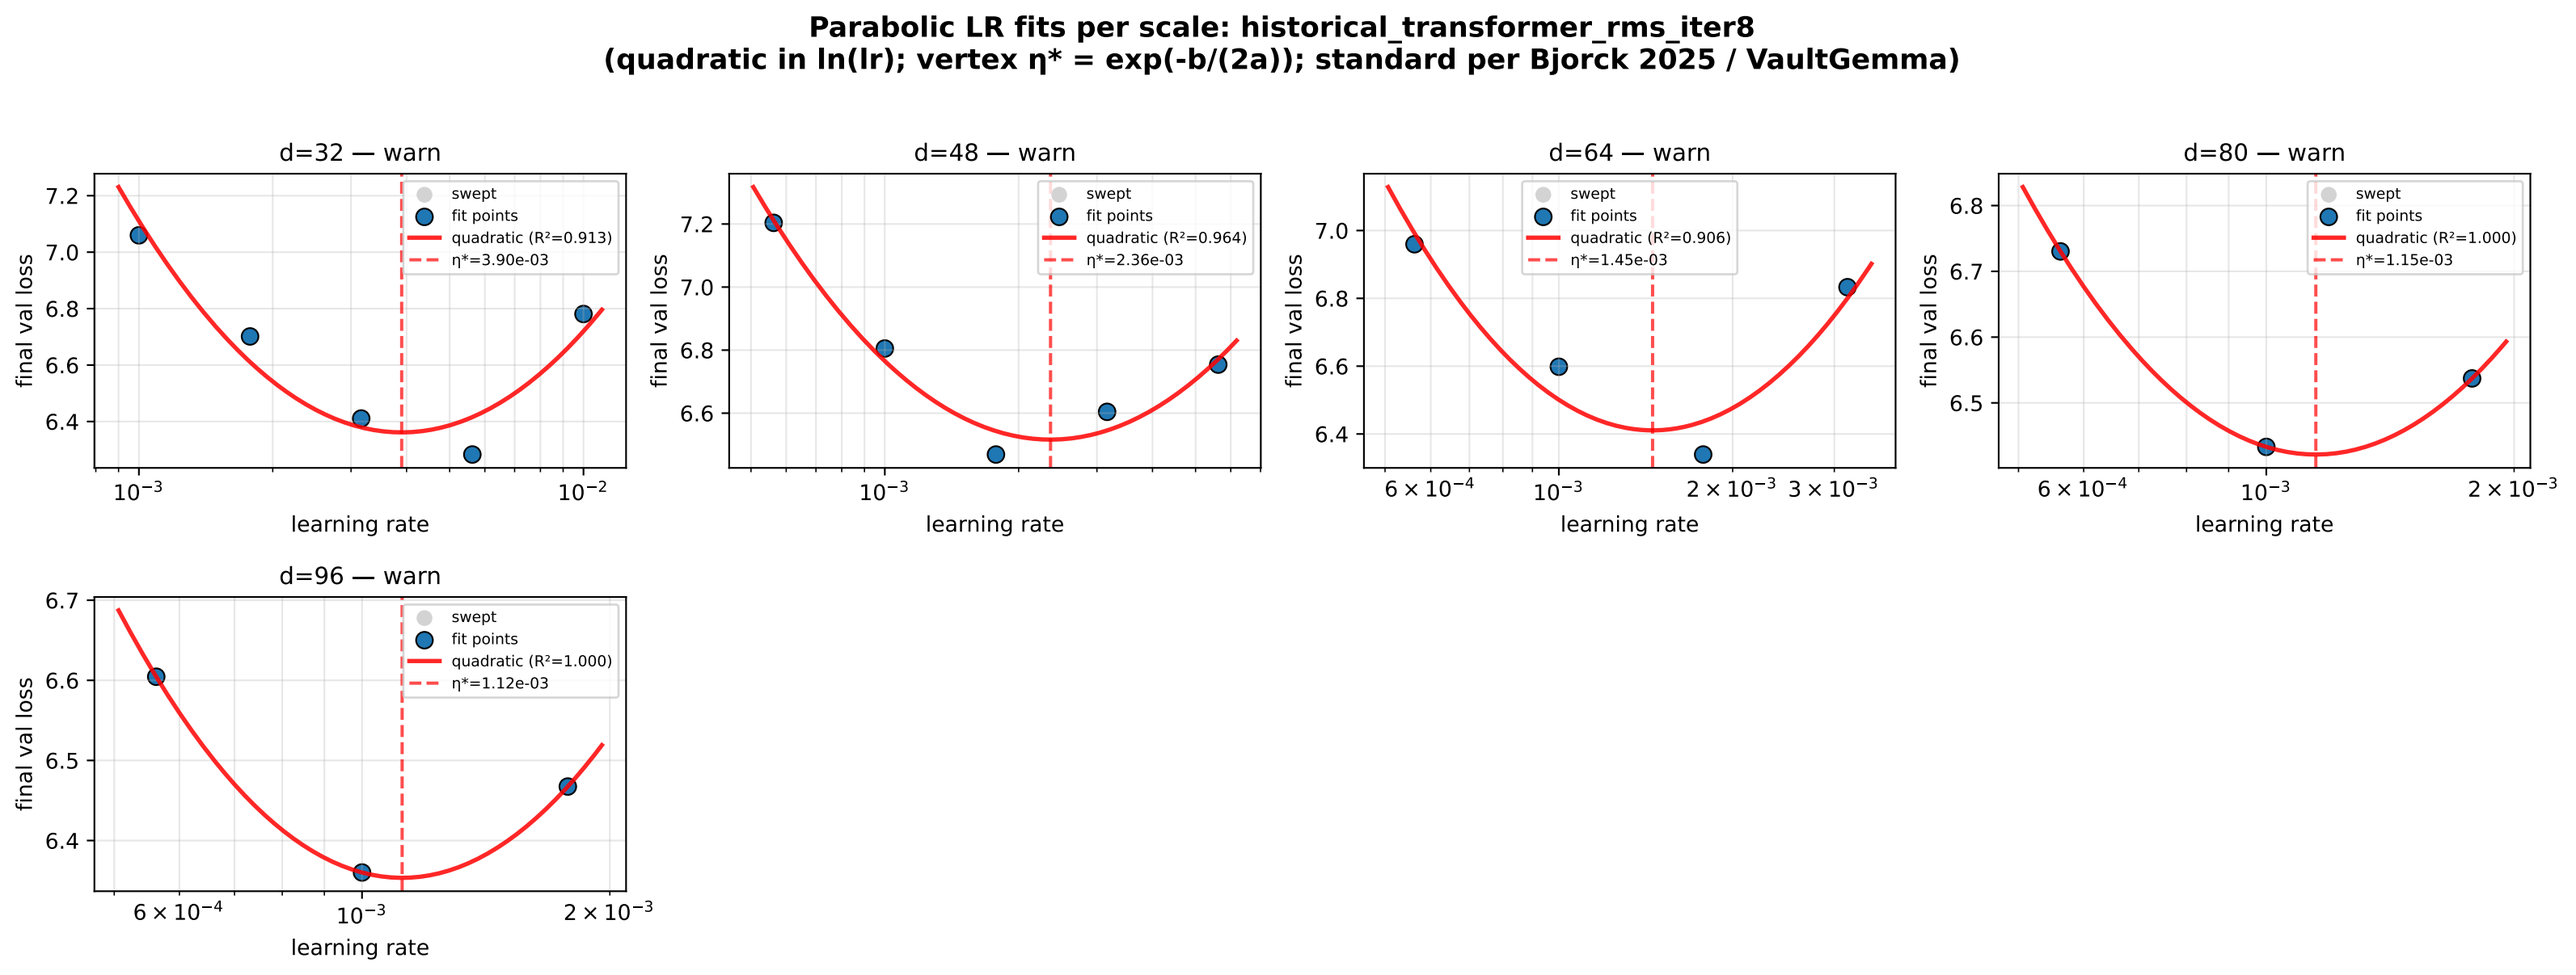
Parabolic LR fits per scale: historical_transformer_rms_iter8
(quadratic in ln(lr); vertex η* = exp(-b/(2a)); standard per Bjorck 2025 / VaultGemma)
d=32 — warn
d=48 — warn
d=64 — warn
d=80 — warn
d=96 — warn
swept
fit points
quadratic (R²=0.913)
η*=3.90e-03
quadratic (R²=0.964)
η*=2.36e-03
quadratic (R²=0.906)
η*=1.45e-03
quadratic (R²=1.000)
η*=1.15e-03
quadratic (R²=1.000)
η*=1.12e-03
final val loss
learning rate