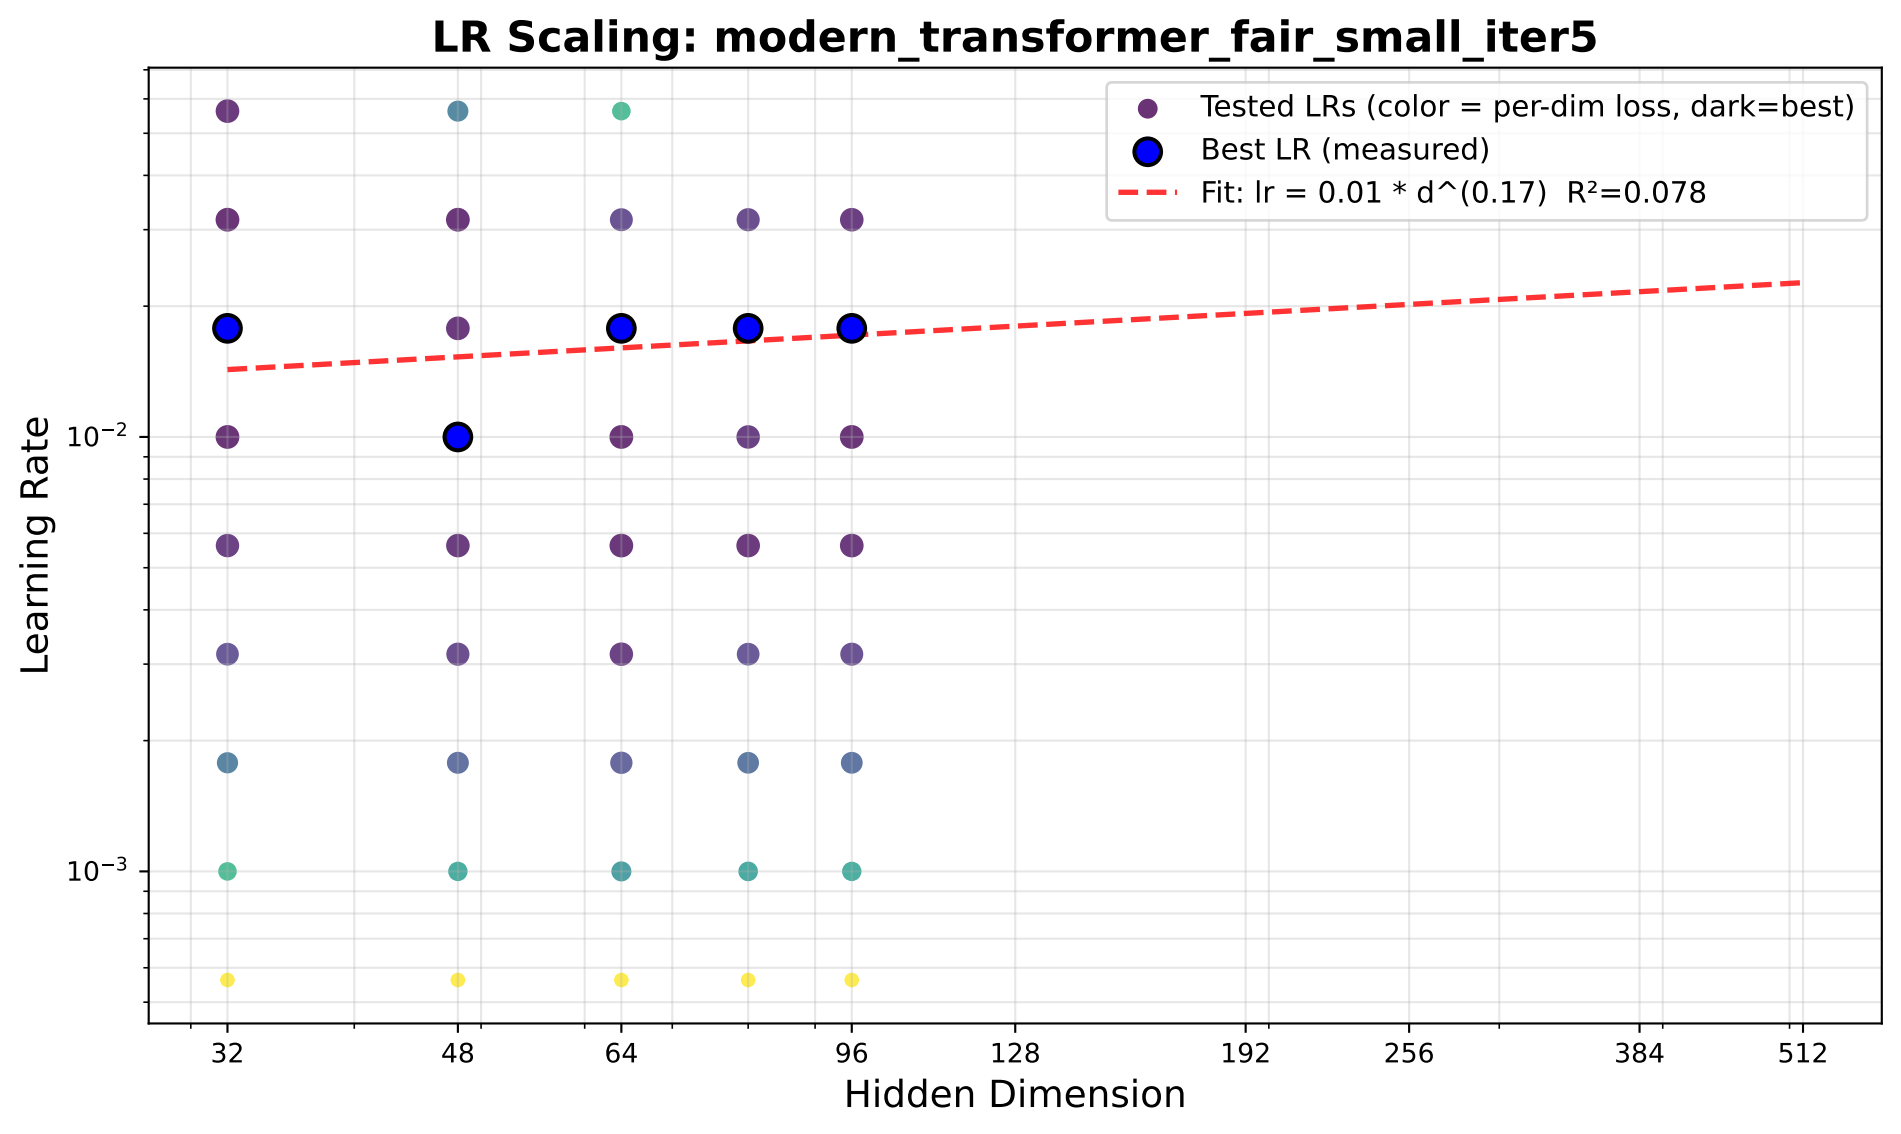
LR Scaling: modern_transformer_fair_small_iter5
Tested LRs (color = per-dim loss, dark=best)
Best LR (measured)
Fit: lr = 0.01 * d^(0.17)  R²=0.078
Learning Rate
$10^{-2}$
$10^{-3}$
32
48
64
96
128
192
256
384
512
Hidden Dimension